

High frequency data with NIO and The DBReader Framework

Basic NIO

- NIO (new input output) is Java's fast data reading / writing library
- It gives us super fast access to raw binary data
- Obviously, speed is very useful to us in the context of the simulation infrastructure we've been talking about

NIO – Reading from channel to buffer

- For example, in the case of a read (code on next slide):
 - We fill a *ByteBuffer* with raw binary data from an open file channel
 - As we fill the buffer, its internal pointer is incremented
 - When we're ready to process data from this buffer, we call its *flip* method to return its internal pointer to the beginning of the buffer
 - We can access that data using convenience methods of *ByteBuffer* such as *getBytes*, *getInt*, *getFloat*, and others
- The same functionality is available for writing

Reading example

```
RandomAccessFile aFile = new RandomAccessFile( "data.txt", "rw" );
FileChannel inChannel = aFile.getChannel();
ByteBuffer buf = ByteBuffer.allocate( 48 );
int bytesRead = inChannel.read( buf );
while ( bytesRead != -1 ) {
    System.out.println( bytesRead );
    buf.flip();
    while(buf.hasRemaining()){
        System.out.print((char) buf.get());
    }
    buf.clear();
    bytesRead = inChannel.read(buf);
}
aFile.close();
```

From *ByteBuffer* documentation

	Absolute <i>get</i> method.
abstract char	<code>getChar()</code> Relative <i>get</i> method for reading a char value.
abstract char	<code>getChar(int index)</code> Absolute <i>get</i> method for reading a char value.
abstract double	<code>getDouble()</code> Relative <i>get</i> method for reading a double value.
abstract double	<code>getDouble(int index)</code> Absolute <i>get</i> method for reading a double value.
abstract float	<code>getFloat()</code> Relative <i>get</i> method for reading a float value.
abstract float	<code>getFloat(int index)</code> Absolute <i>get</i> method for reading a float value.
abstract int	<code>getInt()</code> Relative <i>get</i> method for reading an int value.
abstract int	<code>getInt(int index)</code> Absolute <i>get</i> method for reading an int value.
abstract long	<code>getLong()</code> Relative <i>get</i> method for reading a long value.

Some of the
methods for
interpreting
raw binary
data

The DBReader framework

- In a previous lecture, we described a strategy for reading data from files that may differ in format with time stamps that may not necessarily line up
- That approach is implemented in the DBReader framework, the code for which has been uploaded to the group web site
- The framework is comprised of DBManager and three interface classes – I_DBClock, I_DBReader, and I_DBProcessor

DBManager

- This is the class that manages reading data from potentially many DBReaders
- First the list of readers is sorted by their next time stamp – the time stamp of the next single record that will be read from these readers – so that the reader at the top of the list is chronologically first in line for reading
- The manager then obtains the next time stamp (sequence number) it should read up to from its clock
 - If we want to read every record, the clock would simply return the time stamp obtained from the next reader. However, the clock is useful in reading data in a less granular way – for example, in five minute buckets

DBManager (contd)

- The manager then asks all readers to read all records up to and including the time stamp it obtained from the clock
- When it encounters the first reader that has a time stamp higher than the time stamp obtained from the clock, it stops the reading cycle and starts the processing cycle
- In the processing cycle, iterates through its processors – which, in our case, are simulations – passing to each its list of readers
- The readers are then sorted again, a new time stamp is obtained from the clock, and the loop continues.

Advantages

- DBReader doesn't know anything about the actual file types of the readers
- It works entirely with the interfaces previously mentioned, which are relatively simple
- The readers implement these interfaces so they can be easily incorporated into the DBReader framework
- For the sake of example, I have uploaded the concrete implementations of TAQ quotes and trades DBReaders and related files